

## Session 3: MapReduce Programming on AWS EMR

**Step 1:** Start your AWS EMR instance by logging in to your AWS Management Console.

Clone

Terminate

AWS CLI export

Cluster: upgrad\_emr Waiting Cluster ready after last step completed.

Summary

Application user interfaces

Monitoring

Hardware

Configurations

Events

Steps

Bootstrap actions

### Summary

ID: j-28330VV73D3BB

Creation date: 2021-07-20 19:04 (UTC+5:30)

Elapsed time: 15 minutes

After last step completes: Cluster waits

Termination protection: Off [Change](#)

Tags: -- [View All / Edit](#)

Master public DNS: ec2-54-173-253-106.compute-1.amazonaws.com [Connect to the Master Node Using SSH](#)

### Configuration details

Release label: emr-5.30.1

Hadoop distribution: Amazon 2.8.5

Applications: Hive 2.3.6, Hue 4.6.0, Mahout 0.13.0, Pig 0.17.0, Tez 0.9.2

Log URI: s3://aws-logs-367134191692-us-east-1/elasticmapreduce/ [View](#)

EMRFS consistent view: Disabled

Custom AMI ID: --

### Application user interfaces

Persistent user interfaces [View](#): YARN timeline server, Tez UI

On-cluster user interfaces [View](#): Not Enabled [Enable an SSH Connection](#)

### Network and hardware

Availability zone: us-east-1a

Subnet ID: [subnet-0085edc222d4dad91](#) [View](#)

Master: Running 1 m4.xlarge

Core: --

Task: --

Cluster scaling: Not enabled

### Security and access

Key name: RHEL

EC2 instance profile: EMR\_EC2\_DefaultRole

EMR role: EMR\_DefaultRole

Visible to all users: All [Change](#)

Security groups for Master: [sg-00fcc219431b5c0a6](#) [View](#) (ElasticMapReduce-master)

Security groups for Core & Task: [sg-0a724c5cbae439160](#) [View](#) (ElasticMapReduce-Task: slave)

**Step 2:** Make sure that your instance is up and running fully. Connect to AWS EMR (via PuTTY, etc.). You learnt how to connect to AWS EMR in the previous modules.

### Segment 3: Basic Implementation of MapReduce Using Python

Here, you will learn how to write basic MapReduce scripts for both Mapper and Reducer and simulate a basic MapReduce program using Python. You will be writing a simple MapReduce program to find the maximum age of the customers gender-wise, as discussed in the previous segment.

**Step 1:** First, you need to make a folder in your local EMR instance where you will import your data set and store your scripts for ease of use.

```
[hadoop@ip-172-31-47-91 ~]$ mkdir MapReduce_test
```

**Step 2:** Now, navigate to this folder.

```
[hadoop@ip-172-31-47-91 ~]$ cd MapReduce_test/
```

Now that you have made a directory to work on, you can start importing the data set to the newly created directory of your instance using WinSCP.

After importing the CustomerAge data set to your directory, you can start writing the Mapper script.

**Step 3:** Start by making a new file for your Mapper, say '**mapperage.py**', using vi, and then go into input mode by pressing 'I' on your keyboard.

```
[hadoop@ip-172-31-47-91 MapReduce_test]$ vi mapperage.py
```

**Step 4:** Now input the following code into the file:

```
import sys
# input comes from STDIN (standard input)
for line in sys.stdin:
    line = line.strip()
    line = line.split(",")

    if len(line) >=2:
        gender = line[1]
        age = line[2]
        print ('%s\t%s' % (gender, age))
```

**Step 5:** Save the file by going into command mode by pressing Esc, then type the following, and press Enter:

```
:wq!
```

Now, you will write the Reducer script.

**Step 6:** Make a new file for the Reducer script, say '**reduced\_max.py**', using vi, and go into input mode by pressing 'i' on your keyboard.

```
[hadoop@ip-172-31-47-91 MapReduce_test]$ vi reduced_max.py
```

**Step 7:** Now input the following code into the file:

```
#!/usr/bin/python
#Reducer.py
import sys
gender_age = {}

for line in sys.stdin:
    line = line.strip()
    gender, age = line.split('\t')

    if gender in gender_age:
        gender_age[gender].append(int(age))
    else:
        gender_age[gender] = []
        gender_age[gender].append(int(age))

#Reducer
for gender in gender_age.keys():
    max_age = max(gender_age[gender])
    print(gender + "\t" + str(max_age))
```

Save the file again by typing '**:wq!**' after going into command mode. With this, you have written the scripts for both Mapper and Reducer for your MapReduce program.

**Step 8:** Before running the MapReduce program, you need to give execute permissions to your scripts. Run the following command:

```
[hadoop@ip-172-31-47-91 MapReduce_test]$ chmod 700 mapperage.py  
reduced_max.py
```

**Step 9:** Now, to run the MapReduce job, run the following command:

```
[hadoop@ip-172-31-47-91 MapReduce_test]$ cat CustomerAge.txt | python  
mapperage.py | python reduced_max.py
```

## Segment 4: Hadoop Streaming

In this segment, you will learn how Hadoop Streaming can be used to run your MapReduce job on the Hadoop cluster.

**Step 1:** First, navigate to the test directory that you created in the previous segment, wherein your data set and scripts are stored.

```
[hadoop@ip-172-31-47-91 ~]$ cd MapReduce_test
```

**Step 2:** Now, you need to push your data set into the HDFS cluster before you can run the MapReduce job on it. Run the following command:

```
[hadoop@ip-172-31-47-91 MapReduce_test]$ hadoop fs -put CustomerAge.txt  
/user/hadoop/
```

**Step 3:** Now that you have your data set in the Hadoop cluster, you can run the MapReduce job using Hadoop Streaming. Run the following command:

```
[hadoop@ip-172-31-47-91 MapReduce_test]$ hadoop jar  
/lib/hadoop-mapreduce/hadoop-streaming-2.8.5-amzn-6.jar \  
-file mapperage.py -mapper 'python mapperage.py' \  
-file reduced_max.py -reducer 'python reduced_max.py' \  
-input /user/hadoop/CustomerAge.txt \  
-output /user/hadoop/output_1
```

**Note:** The Hadoop streaming jar that you will be using for your instance may have a different version. You can find out what jar you can use for using Hadoop Streaming by using the 'ls'

command to check out what jar files are present in the '/lib/hadoop-mapreduce/' directory of your instance.

**Note:** You need to always make sure that the output directory that you are specifying in your MapReduce job is not already present in the HDFS cluster; otherwise, the job will throw an error.

**Note:** You can also set the number of Reduce tasks for running your job by setting the argument '**-numReduceTasks**' to a number of your choice. By default, it is set to 1; this should be set to a low number, as a larger number will increase the number of partitions created in the HDFS.

**Step 4:** Now you can check the output of the MapReduce job. Run the following command:

```
[hadoop@ip-172-31-47-91 MapReduce_test]$ hadoop fs -ls  
/user/hadoop/output_1
```

**Step 5:** To check the actual output, type the following command:

```
[hadoop@ip-172-31-47-91 MapReduce_test]$ hadoop fs -cat  
/user/hadoop/output_1/part*
```

## Segment 5: The Combiner

In this segment, you will be calculating the average of the ages of the customers corresponding to their gender with the help of a Combiner.

**Step 1:** Let's first create the Mapper script for your program. Follow the steps mentioned in the previous segment and create a script, say '**mapperage\_avg.py**', and input the following code into the file:

```
import sys  
# input comes from STDIN (standard input)  
for line in sys.stdin:  
    line = line.strip()  
    line = line.split(",")  
  
    if len(line) >= 2:  
        gender = line[1]  
        age = line[2]
```

```
print ('%s\t1\t%s' % (gender, age))
```

**Step 2:** Now you need to write the Combiner script. Create a file named 'combiner\_avg.py' and input the following code into the file:

```
import sys
gender_dict = {}
gender_count, gender_sum = 0, 0

for line in sys.stdin:
    line = line.strip()
    gender, count, age = line.split( '\t' )
    count, age = int(count), int(age)

    if gender in gender_dict:
        gender_dict[gender][1]+=age
        gender_dict[gender][0]+=count
    else :
        gender_dict[gender] = []
        gender_dict[gender].append(count)
        gender_dict[gender].append(age)

#Reducer
for gender in gender_dict.keys():

    print(gender+"\t"+str(gender_dict[gender][0])+"\t"+str(gender_dict[gender][1]))
```

**Note:** Please take care of proper indentation while typing the python code.

**Step 3:** Now you need to write the Reducer for this program. Create a new file, say 'reduced\_avg.py', and input the following code into the file:

```
import sys

gender_dict = {}
gender_count, gender_sum = 0, 0

for line in sys.stdin:
    line = line.strip()
```

```

gender, count, age = line.split('\t')
count, age = int(count), int(age)

if gender in gender_dict:
    gender_dict[gender][1]+=age
    gender_dict[gender][0]+=count
else:
    gender_dict[gender] = []
    gender_dict[gender].append(count)
    gender_dict[gender].append(age)

#Reducer
for gender in gender_dict.keys():
    flag = float(gender_dict[gender][1])/float(gender_dict[gender][0])
    print(gender+"\t"+str(flag))

```

**Note:** Make sure that you give execute permissions to the code files created before running the MapReduce job using the 'chmod 700' command.

**Note:** Here you can try to run these files locally to check if the code is working properly before running it on Hadoop. For this you can use the command - 'cat CustomerAge.csv | python mapperage\_avg.py | python combiner\_avg.py | python reduced\_avg.py'

**Note:** In Hadoop, even if a program fails, the output directory may be created and so always remember to check that the output directory has not already been created before running a MapReduce job.

**Step 4:** Now you can run the MapReduce job to find the average in question. Run the following command:

```

[hadoop@ip-172-31-47-91 MapReduce_test]$ hadoop jar
/lib/hadoop-mapreduce/hadoop-streaming-2.8.5-amzn-6.jar \
-file mapperage_avg.py -mapper 'python mapperage_avg.py' \
-file reduced_avg.py -reducer 'python reduced_avg.py' \
-input /user/hadoop/CustomerAge.txt \
-output /user/hadoop/output_2 \
-file combiner_avg.py -combiner 'python combiner_avg.py'

```

**Step 5:** You can check the output of the program again by typing the following command:

```
[hadoop@ip-172-31-47-91 MapReduce_test]$ hadoop fs -cat  
/user/hadoop/output_2/part*
```

## Segment 6: The Partitioner

In this segment, you will use the airline data set and calculate the average flight delay corresponding to each year in the data set. Since most queries will be based on the year, you need to use the Partitioner so that all key-value pairs belonging to one year are processed by one Reduce task.

**Step 1:** You need to make another directory in your test directory to avoid any confusion between the previous MapReduce scripts and the ones that you will be writing now.

```
[hadoop@ip-172-31-47-91 MapReduce_test]$ mkdir AirlineAverage/
```

```
[hadoop@ip-172-31-47-91 MapReduce_test]$ cd AirlineAverage/
```

**Step 2:** You need to push the airline data set into the HDFS cluster after importing it into the EMR instance by running the 'put' command:

```
[hadoop@ip-172-31-47-91 AirlineAverage]$ hadoop fs -put ../AirLineData.csv  
/user/hadoop/demoAirLine/input/
```

**Note:** Here note that you need to first make a demoAirLine directory and its subdirectories before transferring the AirLineData.csv file to this directory or in MapReduce\_test/ directory, in HDFS. You can do this by using the 'hadoop fs -mkdir /user/hadoop/demoAirLine' and so on for the other subdirectories.

**Step 3:** Now you need to write the Mapper script. Create a file named 'mapper\_airlineAvg.py' and input the following code into the file:

```
import sys  
for line in sys.stdin:  
    line = line.strip()  
    line = line.split(",")  
    if line[0] == "ID":  
        continue
```



```
year = line[1]
arrdel = line[15]
depdel = line[16]
if arrdel == "NA" or depdel == "NA":
    continue
arrdel = int(arrdel)
depdel = int(depdel)
if arrdel+depdel>0:
    total_del = arrdel+depdel
    print("{}\t1\t{}".format(year, total_del))
```

**Step 4:** Now, to write the Combiner script, create a file named '**combiner\_airlineAvg.py**' and input the following code into the file:

```
import sys

flight = {}
flight_count, del_sum = 0, 0

for line in sys.stdin:
    line = line.strip()
    year, count, delay = line.split("\t")
    count, delay = int(count), int(delay)

    if year in flight:
        flight[year][1] += delay
        flight[year][0] += count
    else:
        flight[year] = []
        flight[year].append(count)
        flight[year].append(delay)

for year in flight.keys():
    print("{}\t{}\t{}".format(year, flight[year][0], flight[year][1]))
```

**Step 4:** After this, you need to write the Reducer script. For this, create a file named '**reducer\_airlineAvg.py**' and input the following code into the file:

```
import sys
```

```
flight = {}
flight_count, del_sum = 0, 0

for line in sys.stdin:
    line = line.strip()
    year, count, delay = line.split("\t")
    count, delay = int(count), int(delay)

    if year in flight:
        flight[year][1] += delay
        flight[year][0] += count
    else:
        flight[year] = []
        flight[year].append(count)
        flight[year].append(delay)

for year in flight.keys():
    average = float(flight[year][1])/float(flight[year][0])
    print("{}\t{}".format(year, average))
```

**Note:** Make sure that you give execute permissions to the code files created before running the MapReduce job using the 'chmod 700' command.

**Step 4:** You can now run the MapReduce job to find the average in question. Run the following command:

```
[hadoop@ip-172-31-47-91 AirlineAverage]$ hadoop jar
/lib/hadoop-mapreduce/hadoop-streaming-2.8.5-amzn-6.jar \
-files mapper_airlineAvg.py,Reducer_airlineAvg.py,combiner_airlineAvg.py \
-mapper "python mapper_airlineAvg.py" \
-combiner "python combiner_airlineAvg.py" \
-reducer "python reducer_airlineAvg.py" \
-input /user/hadoop/demoAirLine/input/AirLineData.csv \
-output /user/hadoop/demoAirLine/output_1 \
-partitioner org.apache.hadoop.mapred.lib.HashPartitioner -numReduceTasks 5
```

**Step 5:** Now, let's look at the output of one of the part files made after running the MapReduce job:

```
[hadoop@ip-172-31-47-91 AirlineAverage]$ hadoop fs -cat  
/user/hadoop/demoAirLine/output_1/part-00000
```

As you can see, this part file has the result only for one of the years..

Now, if you check the output part files individually for all the other files, you will see that each of them has the output of a single year, which means that each year had a different partition and, therefore, was assigned different Reduce tasks by the Partitioner.

```
[hadoop@ip-172-31-47-91 AirlineAverage]$ hadoop fs -cat  
/user/hadoop/demoAirLine/output_1/part-00001
```

```
[hadoop@ip-172-31-47-91 AirlineAverage]$ hadoop fs -cat  
/user/hadoop/demoAirLine/output_1/part-00002
```

```
[hadoop@ip-172-31-47-91 AirlineAverage]$ hadoop fs -cat  
/user/hadoop/demoAirLine/output_1/part-00003
```

```
[hadoop@ip-172-31-47-91 AirlineAverage]$ hadoop fs -cat  
/user/hadoop/demoAirLine/output_1/part-00004
```